

# Recent development of Gaussian Process Regression

Qi Wang

January, 2021

## 1 Current directions

### 1.1 Deep Neural Networks as Gaussian Processes

1. Neal [1996] first proved that one-layer neural network is equivalent to Gaussian Processes,
2. and then Lee et al. [2018] proved multi-layer NN is also Gaussian Processes and comes up with the model NNGP referring to a GP which corresponds to a deep, infinitely wide neural network.
3. Then Pang et al. [2019] extends this NNGP model to a more general form and use it to solve partial differential equations (PDEs).

(proof and evaluation refer to papers)

### 1.2 Combine GPs with NN

Friedrich [2020] mapped the equations for the GP's predictive mean and variance onto a specific NN of finite size and train it using standard deep learning techniques.

### 1.3 Large scale of data

1. **Exact GPs** Wang et al. [2019] utilizes multi-GPU parallelization and conjugate gradients method to implement GP with  $10^4 - 10^6$  data points (normal GPR is suitable for thousands of data points).
2. **A review paper** Liu et al. [2020] reviews recent advances for improving the scalability and capability of scalable GPs, e.g., prior approximation (DTC, FITC, etc.) and posterior approximations. He also summarized future research directions:
  - Scalable deep GP
  - Scalable multi-task GP
  - Scalable online GP
  - Scalable GP classification

## 1.4 Rates of Convergence for Sparse Variational Gaussian Process Regression

Burt et al. [2019] wins the best paper award at ICML 2019.

## 2 Deep Neural Networks as Gaussian Processes

### 2.1 Gaussian Processes Regression

For  $(x_i, y_i) \in (X_1, Y_1)$ , with noise  $\epsilon_i \in \mathcal{N}(0, \sigma^2)$  being independent with  $f(x_i)$ , we have

$$y_i = f(x_i) + \epsilon_i.$$

The distribution of Gaussian process  $(Y_1, f(X_2))$  is


$$\begin{pmatrix} Y_1 \\ f(X_2) \end{pmatrix} \sim \mathcal{GP} \left( \begin{pmatrix} \mu_{X_1} \\ \mu_{X_2} \end{pmatrix}, \begin{pmatrix} \Sigma_{X_1 X_1} + \sigma^2 I_{n_1} & \Sigma_{X_1 X_2} \\ \Sigma_{X_2 X_1} & \Sigma_{X_2 X_2} \end{pmatrix} \right).$$

And the mean and covariance of posterior  $f(X_2)|Y_1$  is

$$\begin{aligned} \tilde{\mu} &= \mu_{X_2} + \Sigma_{X_2 X_1} (\Sigma_{X_1 X_1} + \sigma^2 I_{n_1})^{-1} (Y_1 - \mu_{X_1}), \\ \tilde{\Sigma} &= \Sigma_{X_2 X_2} - \Sigma_{X_2 X_1} (\Sigma_{X_1 X_1} + \sigma^2 I_{n_1})^{-1} \Sigma_{X_1 X_2} \end{aligned}$$

We use  $\tilde{\mu}$  as the prediction value in machine learning.

### 2.2 Derivations of NNGP kernel



**Fig. 1.** A fully connected neural net with  $L$  hidden layers. For each unit or neuron in hidden layers, there exist one input  $z_i^{l-1}$  and one output  $x_i^l$  with  $l = 1, 1 \dots, L$ . Layer 0 is the input layer.  $x_i$  is the  $i$ -th component of the input vector  $\mathbf{x}$ , and  $z_i^l$  is the  $i$ -th component of the output vector  $\mathbf{z}$ . The dimensions of input and output spaces are  $d_{in}$  and  $d_{out}$ , respectively. The width for hidden layer  $l$  is  $N_l$ . At the center of each unit is the activation function  $\phi(\cdot)$  that transforms input to the corresponding output. Between two successive layers, the weight  $w_{ij}^l$  for  $l = 0, 1, \dots, L$  denotes the contribution of unit  $j$  in layer  $l$  to unit  $i$  in layer  $l + 1$ . Layer  $L + 1$  is the output layer. The bias  $b_i^l$  is attached to unit  $i$  in layer  $l + 1$  for  $l = 0, 1, \dots, L$ . Note that for clarity most of connecting arrows between layers are omitted.

Figure 1: DNN figure from Pang et al. [2019]

Assume parameters  $w_{ij}^l$  are i.i.d.s for all  $i, j, l$ ; AND  $b_i^l$  are i.i.d.s for all  $i, l$ , and

$$\begin{aligned}\mathbb{E}[w_{ij}^l] &= 0, \quad \mathbb{E}[b_i^l] = 0, \\ \text{Var}[w_{ij}^l] &= \frac{\sigma_w^2}{N_l}, \quad \text{Var}[b_i^l] = \sigma_b^2, \\ \text{thus } \mathbb{E}[(w_{ij}^l)^2] &= \frac{\sigma_w^2}{N_l}, \quad \mathbb{E}[(b_i^l)^2] = \sigma_b^2.\end{aligned}$$

As shown in Figure 1, we have

$$\begin{aligned}\mathbf{x}, & \quad z_j^0(\mathbf{x}) = b_j^0 + \sum_{i=1}^{\text{din}} w_{ij}^0 x_i, & \quad \mathbb{E}[z_j^0(\mathbf{x})] &= 0 \\ x_j^1(\mathbf{x}) = \phi(z_j^0(\mathbf{x})), & \quad z_j^1(\mathbf{x}) = b_j^1 + \sum_{i=1}^{\text{din}} w_{ij}^1 x_i^1(\mathbf{x}), & \quad \mathbb{E}[z_j^1(\mathbf{x})] &= 0 \\ \dots & \quad z_j^{L-1}(\mathbf{x}) = b_j^{L-1} + \sum_{i=1}^{N_{L-1}} w_{ij}^{L-1} x_i^{L-1}(\mathbf{x}), & \quad \mathbb{E}[z_j^{L-1}(\mathbf{x})] &= 0 \\ x_j^L(\mathbf{x}) = \phi(z_j^{L-1}(\mathbf{x})), & \quad z_j^L(\mathbf{x}) = b_j^L + \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}), & \quad \mathbb{E}[z_j^L(\mathbf{x})] &= 0\end{aligned}$$

The expectation of pre-activation variable  $\mathbb{E}[z_j^L(\mathbf{x})]$  is  $\mathbf{0}$  because the parameters  $w_{ij}^l$  and data points  $\mathbf{x} \in \mathbb{R}^{\text{din}}$  are all independent for any  $i, j, l$ . Thus the covariance of data points  $\mathbf{x}$  and  $\mathbf{x}'$  after neural network transformation at Layer  $L$  node  $j$  is (note we shorten  $k^L(z_j^L(\mathbf{x}), z_j^L(\mathbf{x}'))$  to  $k^L(\mathbf{x}, \mathbf{x}')$ )

$$\begin{aligned}k^L(\mathbf{x}, \mathbf{x}') &= k^L(z_j^L(\mathbf{x}), z_j^L(\mathbf{x}')) \\ &= \mathbb{E}[z_j^L(\mathbf{x})z_j^L(\mathbf{x}')] = \mathbb{E}[(b_j^L + \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}))(b_j^L + \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}'))] \\ &= \mathbb{E}[(b_j^L)^2 + b_j^L(\sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}) + \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}')) + \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}) \sum_{i=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}')] \\ &= \sigma_b^2 + 0 + \mathbb{E}[\sum_{i=1}^{N_L} \sum_{k=1}^{N_L} w_{ij}^L x_i^L(\mathbf{x}) w_{kj}^L x_k^L(\mathbf{x}')] \\ &= \sigma_b^2 + \sum_{i=1}^{N_L} \mathbb{E}[(w_{ij}^L)^2] \mathbb{E}[x_i^L(\mathbf{x}) x_i^L(\mathbf{x}')] \\ &= \sigma_b^2 + \frac{\sigma_w^2}{N_L} \sum_{i=1}^{N_L} \mathbb{E}[x_i^L(\mathbf{x}) x_i^L(\mathbf{x}')] \\ &= \sigma_b^2 + \sigma_w^2 \mathbb{E}[x_i^L(\mathbf{x}) x_i^L(\mathbf{x}')] \\ &= \sigma_b^2 + \sigma_w^2 \mathbb{E}[\phi(z_i^{L-1}(\mathbf{x})) \phi(z_i^{L-1}(\mathbf{x}'))] \\ &= \sigma_b^2 + \sigma_w^2 \int \int \phi(z_i^{L-1}(\mathbf{x})) \phi(z_i^{L-1}(\mathbf{x}')) p(z_i^{L-1}(\mathbf{x}), z_i^{L-1}(\mathbf{x}')) d(z_i^{L-1}(\mathbf{x})) d(z_i^{L-1}(\mathbf{x}'))\end{aligned} \tag{1}$$

As we know  $(z_i^{L-1}(\mathbf{x}), z_i^{L-1}(\mathbf{x}'))$  is a Gaussian Processes with distribution of  $\mathcal{GP}(\mathbf{0}, \Sigma)$ , where

$$\Sigma = \begin{pmatrix} k^{L-1}(\mathbf{x}, \mathbf{x}) & k^{L-1}(\mathbf{x}, \mathbf{x}') \\ k^{L-1}(\mathbf{x}', \mathbf{x}) & k^{L-1}(\mathbf{x}', \mathbf{x}') \end{pmatrix} = \begin{pmatrix} k^{L-1}(\mathbf{x}, \mathbf{x}) & k^{L-1}(\mathbf{x}, \mathbf{x}') \\ k^{L-1}(\mathbf{x}, \mathbf{x}') & k^{L-1}(\mathbf{x}', \mathbf{x}') \end{pmatrix}.$$

Thus the bivariate pdf is

$$p(z_i^{L-1}(\mathbf{x}), z_i^{L-1}(\mathbf{x}')) = (2\pi)^{-1} \det(\mathbf{\Sigma})^{-\frac{1}{2}} \exp \left( -\frac{1}{2} \begin{pmatrix} z_i^{L-1}(\mathbf{x}) \\ z_i^{L-1}(\mathbf{x}') \end{pmatrix}^T \mathbf{\Sigma}^{-1} \begin{pmatrix} z_i^{L-1}(\mathbf{x}) \\ z_i^{L-1}(\mathbf{x}') \end{pmatrix} \right) \quad (2)$$

Plug 2 into 1 and calculate the integral, we can write 1 as a recursive function

$$k^L(\mathbf{x}, \mathbf{x}') = \sigma_b^2 + \sigma_w^2 F_\phi(k^{L-1}(\mathbf{x}, \mathbf{x}), k^{L-1}(\mathbf{x}, \mathbf{x}'), k^{L-1}(\mathbf{x}', \mathbf{x}')) \quad (3)$$

For input layer  $L = 0$ , we have

$$\begin{aligned} k^0(\mathbf{x}, \mathbf{x}') &= \mathbb{E}[z_j^0(\mathbf{x}) z_j^0(\mathbf{x}')] = \mathbb{E}[(b_j^0 + \sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}))(b_j^0 + \sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}'))] \\ &= \mathbb{E}[(b_j^0)^2 + b_j^0(\sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}) + \sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}')) + \sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}) \sum_{i=1}^{d_{in}} w_{ij}^0 x_i(\mathbf{x}')] \\ &= \sigma_b^2 + \sum_{i=1}^{d_{in}} \mathbb{E}[(w_{ij}^0)^2] x_i(\mathbf{x}) x_i(\mathbf{x}') \\ &= \sigma_b^2 + \frac{\sigma_w^2}{d_{in}} \mathbf{x}^T \mathbf{x}'. \end{aligned}$$

Thus we can calculate  $k^L(\mathbf{x}, \mathbf{x}')$  recursively from layer 0 to layer  $L$ .

### 2.3 Analytical recursive kernel function

1. When  $\phi$  is **Relu**, Lee et al. [2018] utilized the result of Cho and Saul [2009], writing the analytical form of recursive kernel function for  $l \geq 1$ , where

$$k^l(z_j^l(\mathbf{x}), z_j^l(\mathbf{x}')) \quad (4)$$

$$= \sigma_b^2 + \sigma_w^2 \int_{\mathbb{R}^2} \phi(z_i^{l-1}(\mathbf{x})) \phi(z_i^{l-1}(\mathbf{x}')) p(z_i^{l-1}(\mathbf{x}), z_i^{l-1}(\mathbf{x}')) d(z_i^{l-1}(\mathbf{x})) d(z_i^{l-1}(\mathbf{x}')) \quad (5)$$

$$= \sigma_b^2 + \sigma_w^2 \int_0^\infty \int_0^\infty z_i^{l-1}(\mathbf{x}) z_i^{l-1}(\mathbf{x}') \frac{1}{2\pi \sqrt{\det(\mathbf{\Sigma})}} \exp \left( -\frac{1}{2} \begin{pmatrix} z_i^{l-1}(\mathbf{x}) \\ z_i^{l-1}(\mathbf{x}') \end{pmatrix}^T \mathbf{\Sigma}^{-1} \begin{pmatrix} z_i^{l-1}(\mathbf{x}) \\ z_i^{l-1}(\mathbf{x}') \end{pmatrix} \right) d(z_i^{l-1}(\mathbf{x})) d(z_i^{l-1}(\mathbf{x}')) \quad (6)$$

$$= \sigma_b^2 + \sigma_w^2 \frac{1}{2\pi} \sqrt{k_{11}^{l-1} k_{22}^{l-1}} (\sin \theta + (\pi - \theta) \cos \theta) \quad (7)$$

$$= \sigma_b^2 + \sigma_w^2 \frac{1}{2\pi} \left( \sqrt{\det(\mathbf{\Sigma})} + (\pi - \theta) k_{12}^{l-1} \right), \quad (8)$$

$$\text{where } \theta = \arccos \left( \frac{k_{12}^{l-1}}{\sqrt{k_{11}^{l-1} k_{22}^{l-1}}} \right), \quad \text{and } \mathbf{\Sigma} = \begin{pmatrix} k_{11}^{l-1} & k_{12}^{l-1} \\ k_{12}^{l-1} & k_{22}^{l-1} \end{pmatrix} \quad (9)$$

**TODO**, derive equation (6) to (7) by hand.

2. When  $\phi$  is **Erf**:

$$\text{Erf}(x) = \frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$$

which is similar to  $\tanh(x)$  as shown in Figure 2, the analytical form of kernel function is

$$k^l(z_j^l(\mathbf{x}), z_j^l(\mathbf{x}')) \quad (10)$$

$$= \sigma_b^2 + \sigma_w^2 \int_{\mathbb{R}^2} \phi(z_i^{l-1}(\mathbf{x})) \phi(z_i^{l-1}(\mathbf{x}')) p(z_i^{l-1}(\mathbf{x}), z_i^{l-1}(\mathbf{x}')) d(z_i^{l-1}(\mathbf{x})) d(z_i^{l-1}(\mathbf{x}')) \quad (11)$$

$$= \sigma_b^2 + \sigma_w^2 \frac{2}{\pi} \arctan\left(\frac{2k_{12}^{l-1}}{\sqrt{k_{11}^{l-1}k_{22}^{l-1} - 4(k_{12}^{l-1})^2}}\right) \quad (12)$$

Equation 12 is implemented in package (<https://github.com/brain-research/nngp>). **TODO**, derive equation (11) to (12) by hand.




Figure 2: function  $\text{erf}(x)$  versus  $\tanh(x)$

3. Package (<https://github.com/brain-research/nngp>) also implements analytical kernels for other nonlinearity activation functions, such as sin, LeakyRelu, Abs, etc.

## 2.4 Nonanalytical integration

For some nonlinearity activation function  $\phi$ , such as  $\tanh$ , Lee et al. [2018] proposed an approximation of the double gaussian integration in equation 1 (by populating a table  $F_\phi(i, j)$  in advance and calculate  $F_\phi(\cdot)$  by bilinear interpolation). However, the result is very bad when

I test paper’s open code (<https://github.com/brain-research/nngp>) using my dataset to do a regression task.

## 2.5 NNGP open library

Lee and Novak et al. [2019] continued working on infinite wide neural networks and are developing an open library for it (<https://github.com/google/neural-tangents>) with a documentation paper Novak et al. [2019]. It provides efficient APIs to apply NNGP kernel and NTK kernel. NTK kernel represents the infinite-width neural networks trained by gradient descent. Several papers have shown that randomly initialized neural networks trained with gradient descent are characterized by a distribution that is related to the NNGP, and is described by the so-called Neural Tangent Kernel (NTK).

I also tested my dataset using this library, the result is shown as in Table 1 I am not

Table 1: Test result

nonlinearity	depth	noise_var	nngp			ntk		
			rmse_train	rmse_test	model_time(s)	rmse_train	rmse_test	model_time(s)
relu	1	0	nan	nan	8	nan	nan	12
relu	2	0	nan	nan	12	nan	nan	16
relu	5	0	nan	nan	25	nan	nan	29
relu	10	0	nan	nan	48	nan	nan	53
tanh	1	0	nan	nan	6	nan	nan	8
tanh	2	0	nan	nan	8	nan	nan	13
tanh	5	0	nan	nan	17	nan	nan	28
tanh	10	0	nan	nan	31	nan	nan	53
relu	1	0.001	4.3699527	4.4619255	8	3.2955	4.0827436	9
relu	2	0.001	4.3921933	4.45558	13	nan	nan	14
relu	5	0.001	4.613788	4.538349	27	nan	nan	29
relu	10	0.001	nan	nan	49	nan	nan	53
tanh	1	0.001	3.5829792	4.1464314	7	nan	nan	8
tanh	2	0.001	3.5964766	4.18663	9	nan	nan	13
tanh	5	0.001	3.6501112	4.2361836	18	nan	nan	28
tanh	10	0.001	3.7202923	4.2845206	35	nan	nan	53
relu	1	0.005	4.441606	4.544863	9	3.785048	4.134365	9
relu	2	0.005	4.3993587	4.505067	13	nan	nan	14
relu	5	0.005	4.3559175	4.4584227	27	nan	nan	29
relu	10	0.005	4.3458934	4.4383006	47	nan	nan	53
tanh	1	0.005	3.7566981	4.1049905	7	nan	nan	8
tanh	2	0.005	3.7273495	4.1020055	9	nan	nan	13
tanh	5	0.005	3.70035	4.100665	18	nan	nan	30
tanh	10	0.005	3.6812818	4.0999956	32	nan	nan	53
relu	1	0.01	4.5197663	4.6207676	8	3.91866	4.1767473	9
relu	2	0.01	4.4816885	4.5841255	13	nan	nan	14
relu	5	0.01	4.4295106	4.531488	25	nan	nan	28
relu	10	0.01	4.401551	4.5011983	46	nan	nan	54
tanh	1	0.01	3.8502622	4.124656	6	nan	nan	8
tanh	2	0.01	3.823687	4.117968	9	nan	nan	13
tanh	5	0.01	3.7982786	4.1122155	18	nan	nan	27
tanh	10	0.01	3.7807302	4.109836	32	nan	nan	52
relu	1	0.05	5.1122756	5.178588	8	4.1414957	4.300323	9
relu	2	0.05	5.0991826	5.165244	13	4.0919614	4.262389	14
relu	5	0.05	5.0740542	5.1386285	25	nan	nan	29
relu	10	0.05	5.058492	5.12151	47	nan	nan	54
tanh	1	0.05	4.029341	4.2054377	6	nan	nan	8
tanh	2	0.05	4.0091906	4.193424	9	nan	nan	14
tanh	5	0.05	3.990197	4.183032	17	nan	nan	29
tanh	10	0.05	3.97672	4.176146	32	nan	nan	53

sure about the architecture of this NNGP in this library (how to deal with the intractable integration.)

## 3 Gradient descent trained infinite-width neural network (NTK)

Lee et al. [2019] proved that gradient-based training of wide neural networks with a squared loss produces test set predictions drawn from a Gaussian process with a particular compositional kernel.

## 4 Next

1. Derive the double Gaussian integration for those analytical nonlinearity case.
2. Understand the proof of NTK and the architecture of NTK in package **Neural Tangents**.
3. Consider one direction: tune parameter by optimization on log marginal likelihood.

## References

- Radford M Neal. Priors for infinite networks. In *Bayesian Learning for Neural Networks*, pages 29–53. Springer, 1996.
- Jaehoon Lee, Yasaman Bahri, Roman Novak, Samuel S Schoenholz, Jeffrey Pennington, and Jascha Sohl-Dickstein. Deep neural networks as gaussian processes. In *International Conference on Learning Representations*, 2018.
- Guofei Pang, Liu Yang, and George Em Karniadakis. Neural-net-induced gaussian process regression for function approximation and pde solution. *Journal of Computational Physics*, 384:270–288, 2019.
- Johannes Friedrich. Neuronal gaussian process regression. *Advances in Neural Information Processing Systems*, 33, 2020.
- Ke Wang, Geoff Pleiss, Jacob Gardner, Stephen Tyree, Kilian Q Weinberger, and Andrew Gordon Wilson. Exact gaussian processes on a million data points. In *Advances in Neural Information Processing Systems*, pages 14648–14659, 2019.
- Haitao Liu, Yew-Soon Ong, Xiaobo Shen, and Jianfei Cai. When gaussian process meets big data: A review of scalable gps. *IEEE Transactions on Neural Networks and Learning Systems*, 2020.
- David Burt, Carl Edward Rasmussen, and Mark Van Der Wilk. Rates of convergence for sparse variational gaussian process regression. In *International Conference on Machine Learning*, pages 862–871, 2019.
- Youngmin Cho and Lawrence Saul. Kernel methods for deep learning. *Advances in neural information processing systems*, 22:342–350, 2009.
- Roman Novak, Lechao Xiao, Jiri Hron, Jaehoon Lee, Alexander A Alemi, Jascha Sohl-Dickstein, and Samuel S Schoenholz. Neural tangents: Fast and easy infinite neural networks in python. *arXiv preprint arXiv:1912.02803*, 2019.
- Jaehoon Lee, Lechao Xiao, Samuel Schoenholz, Yasaman Bahri, Roman Novak, Jascha Sohl-Dickstein, and Jeffrey Pennington. Wide neural networks of any depth evolve as linear models under gradient descent. In *Advances in neural information processing systems*, pages 8572–8583, 2019.